

EECS 126 Notes

Japjot Singh

March 31, 2020

Course notes for EECS 126 taken in Spring 2020.

Contents

1	Tuesday, January 21	2
1.1	Fundamentals	2
1.2	Axioms of Probability (Kolmogov)	3
1.3	Fundamental facts about probability	3
1.4	Discrete Probability	3
1.5	Conditional Probability	3
1.6	Product (Multiplication) Rule	4
1.7	Total Probability	4
1.8	Bayes' Theorem	4
2	Thursday, January 23	5
2.1	Announcement	5
2.2	Birthday Paradox	5
2.3	Bayes Rule False Positive Problem	5
2.4	Independence	5
2.4.1	Conditional Independence	6
2.4.2	Independence of a collection of events	6
3	Poisson Process	7
4	Markov Chains	9
4.1	Discrete Time Markov Chains	9
4.1.1	n -Step Transition Probabilities	9
4.2	Classification of States	10
4.2.1	Periodicity	12
4.3	Steady-State Behavior	12
4.4	Birth-Death Process	14
4.5	Absorption Probabilities and Expected Time to Absorption	14
4.6	Recurrence and Transience	14

1 Tuesday, January 21

Understand problem as an "experiment" and then solve it using tools in your skillset: combinatorics, calculus, common sense.

1.1 Fundamentals

Definition 1. Sample Space Ω of an experiment is the set of all outcomes of the experiment.

Example 1.1

Your experiment is 2 fair coins $\Omega = \{HH, HT, TH, TT\}$ these outcomes (base outcomes) are **mutually exclusive (ME)** and **collectively exhaustive (CE)**

Example 1.2

Toss a coin till the first "Heads" $\Omega = \{H, TH, TTH, \dots\}; |\Omega| = \infty$

Example 1.3

Waiting at the bus-stop for next bus $\Omega = (0, T)$

Visual 1 - We have the experiment which produces outcomes, once you have the outcome space the next definition is the definition of events

Definition 2 (Events). Allowable subsets of Ω (collections of outcomes)

Example 1.4

Get at least 1 Head in experiment 1, $\{HH, HT, TH\}$, $p = \frac{3}{4}$

Defining events carefully is the key to tackling many tough problems.

Example 1.5

Ex 2.2: Get an even number of tosses

Example 1.6

Ex 2.3: Waiting time ≤ 5 min

Definition 3 (Probability Space). A **probability space** $(\Omega, \mathcal{F}, \mathcal{P})$ is a mathematical construct to model "experiments" and has 3 components:

1. Ω is the set of all possible outcomes
2. \mathcal{F} set of all events (composition of outcomes), where each event is a set containing 0 or more base outcomes, \emptyset is a **base outcome** where $\mathbb{P}(\emptyset) = 0$. \mathcal{F} is intuitively a powerset (i.e. for the experiment in example 1.1 $\mathcal{F} = \{\emptyset, \{H, H\}, \{H, T\}, \dots\}$).
3. \mathcal{P} is the probability measure which assigns a number in $[0, 1]$ to each event in \mathcal{F} .

Base outcomes must be ME and CE that is when writing out Ω as a collection of all the base outcomes, they should be the most simplified components.

1.2 Axioms of Probability (Kolmogorov)

What properties do we need the probability measure \mathcal{P} to satisfy?

1. $\mathbb{P}(\emptyset) = 0$
2. $\mathbb{P}(\Omega) = 1$, really just a normalization
3. $\mathbb{P}(A_1 \cup A_2 \cup \dots) = \mathbb{P}(A_1) + \mathbb{P}(A_2) + \dots$ for disjoint (ME) events A_1, A_2, \dots

for disjoint $\mathbb{P}(\bigcup_{i=1}^n A_i) = \sum_{i=1}^n \mathbb{P}(A_i)$

1.3 Fundamental facts about probability

1. $\mathbb{P}(A^c) = 1 - \mathbb{P}(A)$
2. $\mathbb{P}(A \cup B) = \mathbb{P}(A) + \mathbb{P}(B) - \mathbb{P}(A \cap B)$ **Vis2 Venn Diagram**
3. Union-bound $\mathbb{P}(A_1 \cup A_2 \cup \dots \cup A_n) \leq \sum_{i=1}^n \mathbb{P}(A_i)$
4. Inclusion-Exclusion, a generalized version of number 2

Theorem 4 (Inclusion-Exclusion)

$$\mathbb{P}(A_1 \cup A_2 \cup \dots \cup A_n) = \sum_{k=1}^n \sum_{1 \leq i_1} \sum_{\leq i_2} \dots \sum_{i_k \leq n} (-1)^{i+1}$$

Proof. here □

1.4 Discrete Probability

$$\mathbb{P}(A) = \sum_{\omega \in A} \mathbb{P}(\omega)$$

In a uniform sample space, all the outcomes are equally likely so then

$$\mathbb{P}(A) = \frac{|A|}{|\Omega|}$$

1.5 Conditional Probability

Similar to events, conditioning on the right event will bail you out of tricky problems.

Definition 5. $\mathbb{P}(A|B) := \mathbb{P}(\text{Event } A \text{ given that Event } B \text{ has occurred})$

Thus for any event A if $\mathbb{P}(B) \neq 0$,

$$\mathbb{P}(A|B) = \frac{\mathbb{P}(A \cap B)}{\mathbb{P}(B)}$$

Example 1.7

Consider 2 six-sided dice. Let A be the event that the first dice rolls is a 6. Let B be the event the sum of the two dice is 7. Then

$$\mathbb{P}(A|B) = \frac{\mathbb{P}(A \cap B)}{\mathbb{P}(B)} = \frac{\mathbb{P}(\{6, 1\})}{\mathbb{P}(\{6, 1\}, \{5, 2\}, \dots, \{1, 6\})}$$

Similarly $\mathbb{P}(A|\text{sum is 11}) = \frac{1}{2}$

When conditioning on B , B becomes to new Ω .

1.6 Product (Multiplication) Rule**1.7 Total Probability****1.8 Bayes' Theorem**

2 Thursday, January 23

2.1 Announcement

Readings B&T ch1 and 2, HW 1 due next wednesday one minute before midnight

2.2 Birthday Paradox

Assuming a group of n individuals whose birth dates are distributed uniformly at random. Given $k = 365$ days in the year what is the probability that at least 2 people in the group share the same birthday. Our sample space is the consists of each possible set of assignments of birth dates to the n students in the class. Since there are 365 possible days for each of the n students in the group $|\Omega| = k^n = 365^n$. Now we can define our event of interest, A , that at least 2 people have the same birthday. Since this is a hard event to work with we can look at the complement A^c the event that no two people share a birth date. We can reach the solution with a counting argument

$$\mathbb{P}(A^c) = \frac{|A^c|}{|\Omega|} = \frac{365 * 364 * \dots * (365 - (n - 1))}{365^n}$$

or with a probabilistic argument using the chain rule

$$\mathbb{P}(A^c) = 1(1 - \frac{1}{k})(1 - \frac{2}{k}) \dots (1 - \frac{n-1}{k})$$

the latter expression can be approximated using Taylor Series which say $e^x \approx 1 + x$ for $|x| \ll 1$.

$$\mathbb{P}(A^c) \approx 1 \cdot e^{-\frac{1}{k}} \cdot e^{-\frac{2}{k}} \dots e^{-\frac{n-1}{k}}$$

thus $\mathbb{P}(A) = 1 - \mathbb{P}(A^c) \approx 1 - e^{-\frac{n^2}{2k}}$

2.3 Bayes Rule False Positive Problem

Supposes there is a new test for a rare disease.

- If a person has the disease, test positive with $p = 0.95$
- If person does not have disease, test negative with $p = 0.95$
- Random person has the disease with $p = 0.001$

Suppose a person tested positive, what is the probability that person has the disease. Let A be the event has disease and B be the event test positive then by applying Bayes Rule directly

$$\mathbb{P}(A|B) = \frac{(0.95)(0.001)}{(0.95)(0.001) + (0.999)(0.05)} = 0.1875$$

the factor heavily contributing to this number is the prior, how rare the disease is in the first place.

2.4 Independence

Definition 6. Two events are independent if the occurrence of one provides **no information** about the occurrence of the other (i.e. $\mathbb{P}(A|B) = \mathbb{P}(A)$).

insert vis1 Independence can also be written as

$$\mathbb{P}(A \cap B) = \mathbb{P}(A)\mathbb{P}(B)$$

Note: Disjoint events are **not** Independent. Events A and B are disjoint if and only if $\mathbb{P}(A \cap B) = 0 \implies \mathbb{P}(A) = 0 \vee \mathbb{P}(B) = 0$. Thus since Base outcomes of a random experiment are disjoint (ME) and have non-zero probabilities they **must be dependent**.

2.4.1 Conditional Independence

$$\mathbb{P}(A \cap B|C) = \mathbb{P}(A|C) \cdot \mathbb{P}(B|C)$$

Note that

- Dependent events can be conditionally independent
- Independent events can be conditionally dependent

Example 2.1

Consider 2 indistinguishable coins: one is two-tailed and the other is two-headed. You pick one of the 2 coins at random and flip it twice.

Let H_i be the event that the i^{th} flip is a Head ($i = 1, 2$). By itself $\mathbb{P}(H_1) = \mathbb{P}(H_2) = \frac{1}{2}$ and $\mathbb{P}(H_2|H_1) = 1 \neq \mathbb{P}(H_2) = 1/2$. Furthermore, $\mathbb{P}(H_1 \cap H_2|A) = \mathbb{P}(H_1|A)\mathbb{P}(H_2|A \cap H_1) = \mathbb{P}(H_1|A)\mathbb{P}(H_2|A)$ which by definition tells us that H_1, H_2 are conditionally independent given A .

2.4.2 Independence of a collection of events

For all possible subsets of your events $A_{1:n}$, each subset must be independent that is

$$\mathbb{P}\left(\bigcap_{i \in S} A_i\right) = \prod_{i \in S} (\mathbb{P}(A_i)), \forall S$$

where S is any subset of the collection of events. Pairwise independence **does not imply** Joint independence of 3 or more events.

3 Poisson Process

If a bus arrives every 15 minutes, with full certainty, then a passenger showing up at a random time will have to wait on average $\frac{1}{16}(0 + 1 + \dots + 14 + 15) = \frac{1}{2}(0 + 15) = 7.5$ minutes. This is because with equal probability you could arrive when the first bus, or when the second bus comes, or any time in between and it is deterministic that each bus arrives every 15 minutes.

Now let's introduce uncertainty. Suppose the arrival of buses is modeled by a Poisson Process with rate $\lambda = \frac{1}{15}$. This means that we expect 1 arrival every $\lambda x = 1 \rightarrow x = \frac{1}{\lambda} = 15$ units of time (minutes). This means that interarrival time $T_i \sim \exp(15)$ and $\mathbb{E}[T_i] = \frac{1}{15}$. That is, a bus arrives every 15 minutes, with some uncertainty.

Furthermore the time-homogeneity ($\mathbb{P}(k, \tau) = \mathbb{P}(k, t)$, $\forall \tau = t$), independence ($N_\tau \sim \text{Poisson}(\lambda\tau)$, specifically N_τ is memoryless, that is for any time $t > 0$ the history of the process after time t is also a Poisson process and is independent from the history of the process until time t this is true in reverse-time direction as well), and small interval properties hold ($\mathbb{P}(0, \tau) = 1 - \lambda\tau$, $\mathbb{P}(1, \tau) = \lambda\tau$, $\mathbb{P}(k, \tau) = \mathcal{O}(\epsilon)$).

We can illustrate this as a limiting case of a discrete setup for better intuition of the properties. Imagine rolling a 15-face die, one face colored green indicating a bus arrival and 14 colored red indicating the absence of an arrival. How often should I roll this dice to simulate the process. The probability of a green face is $\frac{1}{15}$. Let T equal the number of trials until we see a green face. Since each roll of the dice is $\sim \text{Ber}(\frac{1}{15})$, $T \sim \text{Geo}(\frac{1}{15})$ and $\mathbb{E}[T] = 15$. We expect to roll 15 times to observe one success. How often should we roll to reach our target success rate of $\lambda = \frac{1}{15} = \frac{1 \text{ success}}{15 \text{ minutes}}$ (we can replace minutes with any units but in the question we were given that the rate was defined in terms of minutes). So we just need to solve

$$\frac{1 \text{ success}}{15 \text{ rolls}} \cdot \frac{1 \text{ roll}}{x \text{ minutes}} = \frac{1 \text{ success}}{15 \text{ minutes}} = \lambda$$

$$\rightarrow x = 1$$

If we roll once every minute, then we will expect to see one arrival every 15 minutes. Now let us consider a 60 sided dice with one face green and the rest red. Then $T \sim \text{Geo}(\frac{1}{60})$ and $\mathbb{E}[T] = 60$ implying $\frac{1}{60} \frac{\text{success}}{\text{rolls}}$ and thus

$$\frac{1 \text{ arrival}}{60 \text{ rolls}} \cdot \frac{1 \text{ roll}}{x \text{ minutes}} = \frac{1 \text{ arrival}}{15 \text{ minutes}}$$

$$\implies x = \frac{1}{4}$$

If we roll once every 15 seconds, 0.25 minutes, then we will expect to see one arrival every 15 minutes.

Now we can consider a N sided dice. $T \sim \text{Geo}(\frac{1}{N})$ and $\mathbb{E}[T] = N$, so we expect one arrival every m units (minutes). We need $Nm = 15$ and so $m = \frac{15}{N}$ verified below

$$\frac{1 \text{ arrivals}}{N \text{ trials}} \cdot \frac{1 \text{ trials}}{x \text{ units}} = \frac{1 \text{ arrivals}}{15 \text{ units}}$$

$$\rightarrow x = \frac{15}{N}$$

$$\lim_{N \rightarrow \infty} \frac{1}{\lambda N} \rightarrow 0$$

That is as we let this dice contain an infinite number of sides, the probability of an arrival at an instant goes to 0, but also the rate at which we sample for arrivals gets faster and faster and

we expect to see one arrival every $\frac{1}{\lambda}$ units. This is the continuous time Poisson process.

Notice how the properties of continuous probability and PP line up. **is there a correspondance with density, probability at a single point being 0 and then delta getting smaller corresponding with the interval becoming smaller?** The PP properties also hold and rather intuitively. Since each roll of the dice is a Bernoulli trial, each sample is iid. Independence follows directly from the fact that the number of arrivals during an interval is the sum of the iid samples in that interval and since functions of independent variables are independent, the number of arrivals in disjoint intervals are independent. And since they are identically distributed, time-homogeneity holds. $\mathbb{P}(1, \tau)$ can be written as some quantity times the number of intervals of unit-length m . In a length τ interval there are $\frac{\tau}{m}$ samples from arrivals. In each of these samples the probability of arrival is some a . We want to find a such that $a \cdot \frac{\tau}{m} = \lambda\tau$. We find that such a satisfies $\frac{m}{a} = \frac{1}{\lambda}$ and we use the fact that $Nm = \frac{1}{\lambda}$ from the bottom of the previous page along with $\lim_{N \rightarrow \infty} m = 0$ to get $a = \frac{1}{N}$ and the interpretation that in the limit as $N \rightarrow \infty$ we approach an infinite number samples from arrivals where each sample is an arrival with probability $\frac{1}{N}$. That is in a PP the probability of one arrival in a length τ interval is equivalent to integrating an infinitely small uniform density across a relatively large region with respect to the rate parameter that the proportion of arrivals to unit-length intervals is equal to λ . The probability of more than one arrival in an interval of length τ also approach 0 as τ approaches 0.

Thus our intuition: A Poisson process with rate λ models arrivals over a period of time where the rate indicates how many unit-length intervals δ there are for each arrival. So for some time interval τ . The probability of k arrivals in this interval τ , is effectively equal to the probability of k successful realizations of

$$\lim_{N \rightarrow \infty} \frac{\tau}{\delta} \text{Bernoulli}(1/N)$$

In cleaner notation: $X_i \sim \text{Bernoulli}(\frac{1}{N})$ models a arrival or non-arrival in each sample. We sample a total of $\frac{\tau}{\delta}$ times, or the total number of unit interval times since X_i models arrivals in unit intervals. Furthermore we are modeling $N\delta = \frac{1}{\lambda}$ we have

$$\lim_{N \rightarrow \infty} \delta = 0$$

an inverse relationship between N and δ . Then the number of arrivals of the interval of length τ is just $S_\tau = X_1 + X_2 + \dots + X_{\frac{\tau}{\delta}}$. And so $\mathbb{P}(k, \tau) = \mathbb{P}(S_\tau = k) \sim \text{Binomial}(\frac{\tau}{\delta}, \frac{1}{N})$ and by the Poisson approximation of the Binomial for small p and large n $\mathbb{P}(k, \tau) \text{Poisson}(np) = \text{Poisson}(\frac{\tau}{N\delta})$. Notice that since $N\delta = \frac{1}{\lambda}$ we have $\tau \cdot \frac{1}{N\delta} = \lambda\tau$ thus the number of arrivals in an interval of length τ is $\text{Poisson}(\lambda\tau)$.

Even cooler yet, the memoryless property of Poisson reads off directly from the iid-ness from Bernoulli RVs. Given we have already seen t arrivals, the probability of seeing $s + t$ arrivals is the same as seeing s arrivals. This follows directly from the poisson. But also not that since the Bernoulli's are iid each of the additional s success are not in any way affected by the history of t arrivals. The waiting time until the next arrival is just the waiting time for a Bernoulli to output 1, but that is just geometric and the continuous analog is exponential. Given that we have k arrivals in an interval of length τ its easy to see that the distribution of arrivals is uniform. Think of the discretized version. If we know that in τ iid bernoulli RVs there were k successes and nothing else, then the distribution on the indices of the bernoulli's which were successful is uniform on $[1, \tau]$.

4 Markov Chains

Interested in models where the effect of the past on the future is summarized by a state, which changes over time given probabilities.

4.1 Discrete Time Markov Chains

In **discrete-time Markov chains**, state changes at certain discrete time instants, indexed by an integer variable n . At each step n , the state of the chain is denoted X_n , and belongs to a **finite** set \mathcal{S} of possible states, called the state space. WLOG let $\mathcal{S} = \{1, \dots, m\}$. The Markov Chain is described in terms of its transition probabilities p_{ij} where

$$p_{ij} = \mathbb{P}(X_{n+1} = j | X_n = i) \quad i, j \in \mathcal{S}$$

The key assumption underlying these chains is that the transition probabilities apply whenever i is visited, no matter what happened in the past, and no matter how i was reached, formally this is the **Markov property**, requiring that:

$$\mathbb{P}(X_{n+1} = j | X_n = i, X_{n-1} = i_{n-1}, \dots, X_0 = i_0) = \mathbb{P}(X_{n+1} = j | X_n = i) = p_{ij}$$

Furthermore the transition probabilities p_{ij} must be nonnegative, and sum to one:

$$\sum_{j=1}^m p_{ij} = 1, \text{ for all } i$$

Another efficient way to encode the MC chain model is a transition probability matrix, a 2D array whose element at row i and column j is p_{ij} , the transition probability from i to j

$$P = \begin{bmatrix} p_{11} & p_{12} & \cdots & p_{1m} \\ p_{21} & p_{22} & \cdots & p_{2m} \\ \vdots & \vdots & \ddots & \vdots \\ p_{m1} & p_{m2} & \cdots & p_{mm} \end{bmatrix}$$

Note that the transition matrix format is (row, col), (i, j) , (from, to).

There is often a need to introduce new states that capture the dependence of the future on the model's past history. The probability of any single path can be found simply using the multiplication rule and tracing the path of transition probabilities. If there is no conditioning on the first state then we need to specify a probability law for the initial state X_0 , the initial distribution.

4.1.1 n -Step Transition Probabilities

Many problems require calculating the probability law of the state at some future time, conditioned on the current state. This probability law is captured by the **n -step transition probabilities**, defined by

$$r_{ij}(n) = \mathbb{P}(X_n = j | X_0 = i)$$

In words, $r_{ij}(n)$ is the probability that the state after n time periods will be j , given that the current state is i . We can calculate it using the following recursion, **Chapman-Kolmogorov equation**

Theorem 7 (Chapman-Kolmogorov Equation for n -Step Transition Probabilities)

The n -step transition probabilities can be calculated using the formula

$$r_{ij}(n) = \sum_{k=1}^m r_{ik}(n-1)p_{kj}, \quad n > 1, \text{ and all } i, j$$

starting with $r_{ij}(1) = p_{ij}$

Proof.

$$\begin{aligned} \mathbb{P}(X_n = j | X_0 = i) &= \sum_{k=1}^m \mathbb{P}(X_{n-1} = k | X_0 = i) \mathbb{P}(X_n = j | X_{n-1} = k, X_0 = i) \\ &= \sum_{k=1}^m r_{ik}(n-1) \end{aligned}$$

□

The Chapman-Kolmogorov equation can be represented more concisely via matrix multiplication, specifically the matrix of n -step transition probabilities $r_{ij}(n)$ is obtained by multiplying the matrix of $(n-1)$ -step transition probabilities $r_{ik}(n-1)$, with the one-step transition probability matrix. Thus the n -step transition probability matrix is just the n th power of the transition probability matrix, P^n .

As $n \rightarrow \infty$, if $r_{ij}(n)$ converges to a limit and this limit does not depend on initial state i , then we say that state j has a positive "steady-state" probability of being occupied at times far into the future. However, there are examples of qualitatively different behavior: where $r_{ij}(n)$ converges, but the limit depends on the initial state, and can be zero for selected states, particularly the probability that a particular absorbing state will be reached depends on how "close" we start to that state. This illustrates that there is a variety of states and asymptotic occupancy behavior in Markov chains. Thus we are motivated to classify and analyze various possibilities.

4.2 Classification of States

We begin by focusing on the mechanism by which some states after being visited once, are certain to be visited again, while for other states this may not be the case. Our goal is to classify the states of a Markov chain with a focus on the long-term frequency by which they are visited. Let us first make the notion of revisiting a state precise. A state j is **accessible** from state i if for some n , the n -step transition probability $r_{ij}(n)$ is positive (there is a positive probability of reaching j , starting from i , after some number of time steps). Let $A(i)$ denote the set of states accessible from state i .

Definition 8 (recurrent). State i is **recurrent** if for every j that is accessible from i , i is also accessible from j , that is $\forall j \in A(i), i \in A(j)$.

If we start at a recurrent state i , we can only visit $j \in A(i)$. But since state i is recurrent $i \in A(j)$. Thus, from any future state, there always some probability of returning to state i . Given enough time, this is certain to happen. By repeating this argument, if a recurrent state is visited once, it is certain to be revisited an infinite number of times. **does this hold only if state j is itself recurrent? like what if you transition to state j but then from state j you can transition into some absorbing state k , then there is a nonzero probability of not returning to state i , which means that you are not certain to revisit i an infinite number of times.**

A state is called **transient** if it is not recurrent. Thus, a state i is transient if there is a state $j \in A(i)$ such that i is not accessible from j , that is there exists $j \in A(i)$ but $i \notin A(j)$. After each visit to state i , there is a positive probability that the state enters such a state j from which i is no longer accessible. Given enough time, this will happen, and state i cannot be visited after that. Thus a transient state will only be visited a finite number of times.

Definition 9 (recurrent class). If i is a recurrent state, the set of states $A(i)$ that are accessible from i form a **recurrent class** (or simply **class**). This means that the states in $A(i)$ are all accessible from each other, and no state outside $A(i)$ is accessible from them. Mathematically, for a recurrent state i , we have $A(i) = A(j)$ for all j that belong to $A(i)$.

At least one recurrent state must be accessible from any transient state, this follows from the definition of transient state. It follows that there must exist at least one recurrent state and hence at least one class, giving the following result

Theorem 10 (Markov Chain Decomposition) • A MC can be decomposed into one or more recurrent classes, plus possibly some transient states

- A recurrent state is accessible for all states in its class, but is not accessible from recurrent states in *other classes*
- A transient state is not accessible from any recurrent state, but a recurrent state must be accessible from a transient state (otherwise the state cannot be transient it would just be recurrent with itself)
- At least one, possibly more, recurrent states are accessible from a given transient state, this follows directly from the previous bullet

The previous theorem implies the following:

1. once the state enters (or starts in) a class of recurrent states, it stays within that class; since all states in the recurrence class are accessible from each other, all states in the class will be visited an infinite number of times
2. if the initial state is transient, then the state trajectory contains an initial portion consisting of transient states and a final portion consisting of recurrent states from the same class

Example 4.1

In a MC at least one recurrent state must be accessible from any given state. That is, for any i , there is at least one recurrent j in the set $A(i)$.

Proof. **TODO**

□

Example 4.2

Show that if a recurrent state is visited once, the probability that it will be visited again in the future is equal to 1 (and, therefore, the probability that it will be visited an infinite number of times is equal to 1).

Proof. Let s be a recurrent state, and suppose s has been visited once. From then on, the only possible states are those in the same recurrence class as s . Therefore WLOG we can assume there is a single self-absorbing state s . We now want to show from some current state $i \neq s$, s is guaranteed to be visited some time in the future.

Consider a new MC with a self-absorbing state s . The transitions out of states i , $i \neq s$ are unaffected. Clearly, s is recurrent in the new chain. Furthermore, for any $i \neq s$, there is a positive probability path from i to s in the original chain (since s is recurrent in the original chain), this holds true in the new chain **still not completely convinced of this**. Since i is not accessible from s in the new chain, it follows that every $i \neq s$ in the new chain is transient. But if there exist recurrent states in a MC, they will eventually be visited, so the state s will eventually be visited by the new chain (with probability 1). But the original chain is identical to the new one until the time that s is first visited. Hence, state s is guaranteed to be eventually visited by the original chain. By repeating this argument, we see that s is guaranteed to be visited an infinite number of times (with probability 1). \square

4.2.1 Periodicity

We are now interested in a characterization of a recurrent class, which relates to the presence or absence of a certain periodic pattern in the times that a state can be visited.

Definition 11 (periodic). A recurrent class is said to be periodic if the states can be grouped in $d > 1$ disjoint subsets S_1, \dots, S_d so that all transitions from one subset lead to the next subset. A recurrent class that is not periodic is said to be aperiodic.

In a periodic recurrent class we move through the sequence of subsets in order, and after d steps, we end up in the same subset. Given a periodic recurrent class and a time n , and a state i , there must exist one or more states j for which $r_{ij}(n) = 0$. This is because starting from i , only one of the sets S_k is a possible transition at time n . Thus to verify aperiodicity of a recurrent class R , is to check whether there is a time $n \geq 1$ and a state $i \in R$ from which all states $j \in R$ can be reached in n steps, that is $r_{ij}(n) > 0$, $\forall j \in R$. A converse is true as well: if a recurrent class R is aperiodic, then there exists a time n such that $r_{ij}(n) > 0$, $\forall i, j \in R$. Another way to check periodicity within a recurrent class is to calculate the gcd of all length paths from a state back to itself. If the gcd is 1 then the class is aperiodic. Equivalently, if there is any state with a self-loop in the recurrent class, then the recurrence class must be aperiodic.

4.3 Steady-State Behavior

We wish to understand long-term state occupancy behavior, the n -step transition probabilities $r_{ij}(n)$ when n is very large. We are interested in understanding when $r_{ij}(n)$ converges to steady-state values that are independent of the initial state. If there are multiple recurrence classes, then the limiting values of $r_{ij}(n)$ depend on the initial state. Thus we restrict our attention to chains involving a single recurrent class, plus some transient states and we can generalize the asymptotic behavior of multiclass chains in terms of the asymptotic behavior of single-class chains.

A single recurrent class is not strong enough to guarantee convergence of $r_{ij}(n)$, consider a recurrence class with two states 1 and 2 where $p_{12} = p_{21} = 1$ observe $r_{ij}(n)$ generically oscillates.

Thus if we exclude multiple recurrent classes and/or a periodic class, for every state j the probability $r_{ij}(n)$, of being at state j , approaches a limiting value π_j that is independent of the initial state i with the following interpretation

$$\pi_j \approx \mathbb{P}(X_n = j), \text{ when } n \text{ is large}$$

and is called the **steady-state probability of j** . This is consolidated in the following important theorem

Theorem 12 (Steady-State Convergence Theorem)

Consider a MC with a single recurrent class, which is periodic. Then, the states j are associated with steady-state probabilities π_j that have the following properties:

(a) For each j , we have

$$\lim_{n \rightarrow \infty} r_{ij}(n) = \pi_j, \quad \forall i$$

(b) The π_j are the unique solution to the system of equations below:

$$\begin{aligned} \pi_j &= \sum_{k=1}^m \pi_k p_{kj}, \quad j = 1, \dots, m, \\ 1 &= \sum_{k=1}^m \pi_k \end{aligned}$$

(c) We have

$$\begin{aligned} \pi_j &= 0, \text{ for all transient states } j, \\ \pi_j &> 0, \text{ for all recurrent states } j \end{aligned}$$

The steady-state probabilities π_j sum to 1 and form a probability-distribution π on the state space, called the **stationary distribution** of the chain with the property that if the initial state is chosen according to the distribution π then

$$\pi P^n = \pi, \quad \forall n \in \mathbb{Z}_{\geq 0}$$

Thus if the initial state is chosen according to the stationary distribution, the state at any future time will have the same distribution. The equations

$$\pi_j = \sum_{k=1}^m \pi_k p_{kj}, \quad j = 1, \dots, m,$$

are called the **balance equations**.

Theorem 13 (Steady-State Probabilities as Expected State Frequencies)

For a MC with a single class which is aperiodic, the steady-state probabilities π_j satisfy

$$\pi_j = \lim_{n \rightarrow \infty} \frac{v_{ij}(n)}{n}$$

where $v_{ij}(n)$ is the expected value of the number of visits to state j within the first n transitions, starting from state i .

Theorem 14 (Expected Frequency of a Particular Transition)

Consider n transitions of a MC with a single class which is aperiodic, starting from an initial state. Let $q_{jk}(n)$ be the expected number of such transitions that take the state from j to k . Then, regardless of the initial state we have

$$\lim_{n \rightarrow \infty} \frac{q_{jk}(n)}{n} = \pi_j p_{jk}$$

That is, the expected frequency of transitions from j to k , regardless of initial state converges to $\pi_j p_{jk}$.

4.4 Birth-Death Process

Definition 15 (birth-death process). A birth-death process is a MC in which the states are linearly arranged and transitions can only occur to a neighboring state, or else leave the state unchanged (self-loop).

Birth-death processes also induce the notation

$$b_i = \mathbb{P}(X_{n+1} = i + 1 | X_n = i), \quad \text{"birth" probability at state } i$$

$$d_i = \mathbb{P}(X_{n+1} = i - 1 | X_n = i), \quad \text{"death" probability at state } i$$

For a b-d process the balance equations are substantially simplified. Any transition from i to $i + 1$ has to be followed by a transition from $i + 1$ to i before another transition from i to $i + 1$ to occur. Thus the expected frequency of transitions from i to $i + 1$, which is $\pi_i b_i$ must be equal to the expected frequency of transitions from $i + 1$ to i which is $\pi_{i+1} d_{i+1}$, leading us to the **local balance** equations

$$\pi_i b_i = \pi_{i+1} d_{i+1}, \quad i = 0, 1, \dots, m - 1$$

Using these local balance equations we obtain

$$\pi_i = \pi_0 \frac{b_0 b_1 \cdots b_{i-1}}{d_1 d_2 \cdots d_i}, \quad i = 1, \dots, m$$

using this and the normalization equation we can compute the probabilities π_i .

4.5 Absorption Probabilities and Expected Time to Absorption**4.6 Recurrence and Transience**

For each $x \in \mathcal{X}$, the random variable T_x represents the first time that the chain visits state x , that is $T_x := \min\{n \in \mathbb{N} : X_n = x\}$. T_x is called the hitting time of state x . We further define T_x^+ which is the hitting time for state x except T_x^+ cannot equal 0, to avoid the trivial condition in which the chain starts at x . Let $\rho_{x,y} := \mathbb{P}_x(T_y^+ < \infty)$, the probability that starting from state x we eventually reach state y and shorthand $\rho_x = \rho_{x,x}$.

Definition 16. A state x is **recurrent** if $\rho_x = 1$ and **transient** if $\rho_x < 1$.

Proposition 17. Let N_x be the total number of visits to state x . If x is recurrent then $N_x = \infty$ in probability almost surely. Thus $\mathbb{E}_x[N_x] = \mathbb{E}[N_x | X_0 = x] = \infty$. On the other hand, if x is transient then $\mathbb{E}_x[N_x] < \infty$, in fact,

$$\mathbb{E}_x[N_x] = \frac{\rho_x}{1 - \rho_x}$$